

credit-score-classification

August 22, 2023

1 About Dataset

This dataset contains information about a sample of over 100 people across the world. The data includes the following informations:

Age: The age of the person in years.

Gender: The gender of the person (male or female).

Income: The annual income of the person in US dollars.

Education: The highest level of education completed by the person.

Marital Status: The marital status of the person (single, married, divorced, or widowed).

Number of Children: The number of children the person has.

Home Ownership: Whether the person owns their home or rents it.

Credit Score: The credit score of the person, which is a measure of their creditworthiness.

1.1 Source and link

Source: Kaggle

Link: <https://www.kaggle.com/datasets/sujithmandala/credit-score-classification-dataset>

Author: Bao Thai

2 Credit Score Classification

2.1 Import needed libraries and packages

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score, \
    classification_report
from sklearn.linear_model import LogisticRegression
```

```

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from xgboost import XGBClassifier

```

2.2 Loading Dataset

```

[2]: df = pd.read_csv('C:/Users/Admin/Desktop/Projects/Credit Score Classification/
↳Credit Score Classification Dataset.csv')
df_original = df.copy()
df_original.head(10)

```

```

[2]:   Age  Gender  Income      Education Marital Status \
0    25  Female   50000  Bachelor's Degree      Single
1    30   Male  100000   Master's Degree    Married
2    35  Female   75000      Doctorate    Married
3    40   Male  125000  High School Diploma      Single
4    45  Female  100000  Bachelor's Degree    Married
5    50   Male  150000   Master's Degree    Married
6    26  Female   40000  Associate's Degree      Single
7    31   Male   60000  Bachelor's Degree      Single
8    36  Female   80000   Master's Degree    Married
9    41   Male  105000      Doctorate      Single

      Number of Children Home Ownership Credit Score
0                0      Rented      High
1                2      Owned      High
2                1      Owned      High
3                0      Owned      High
4                3      Owned      High
5                0      Owned      High
6                0      Rented    Average
7                0      Rented    Average
8                2      Owned      High
9                0      Owned      High

```

2.3 Explore Data

```

[16]: df_original.shape

```

```

[16]: (164, 8)

```

```

[17]: df_original.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 164 entries, 0 to 163
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                   164 non-null   int64
1   Gender                164 non-null   object
2   Income                164 non-null   int64
3   Education             164 non-null   object
4   Marital Status        164 non-null   object
5   Number of Children    164 non-null   int64
6   Home Ownership        164 non-null   object
7   Credit Score          164 non-null   object
dtypes: int64(3), object(5)
memory usage: 10.4+ KB

```

```
[18]: df_original.isna().sum()
```

```

[18]: Age                0
      Gender            0
      Income            0
      Education         0
      Marital Status    0
      Number of Children 0
      Home Ownership    0
      Credit Score      0
      dtype: int64

```

```
[19]: df_original.describe()
```

```

[19]:
count    Age      Income  Number of Children
count  164.000000   164.000000         164.000000
mean    37.975610  83765.243902          0.652439
std      8.477289  32457.306728          0.883346
min     25.000000  25000.000000          0.000000
25%     30.750000  57500.000000          0.000000
50%     37.000000  83750.000000          0.000000
75%     45.000000 105000.000000          1.000000
max     53.000000 162500.000000          3.000000

```

```
[20]: df_original.corr()
```

```

C:\Users\Admin\AppData\Local\Temp\ipykernel_10520\2946492221.py:1:
FutureWarning: The default value of numeric_only in DataFrame.corr is
deprecated. In a future version, it will default to False. Select only valid
columns or specify the value of numeric_only to silence this warning.
df_original.corr()

```

```
[20]:
```

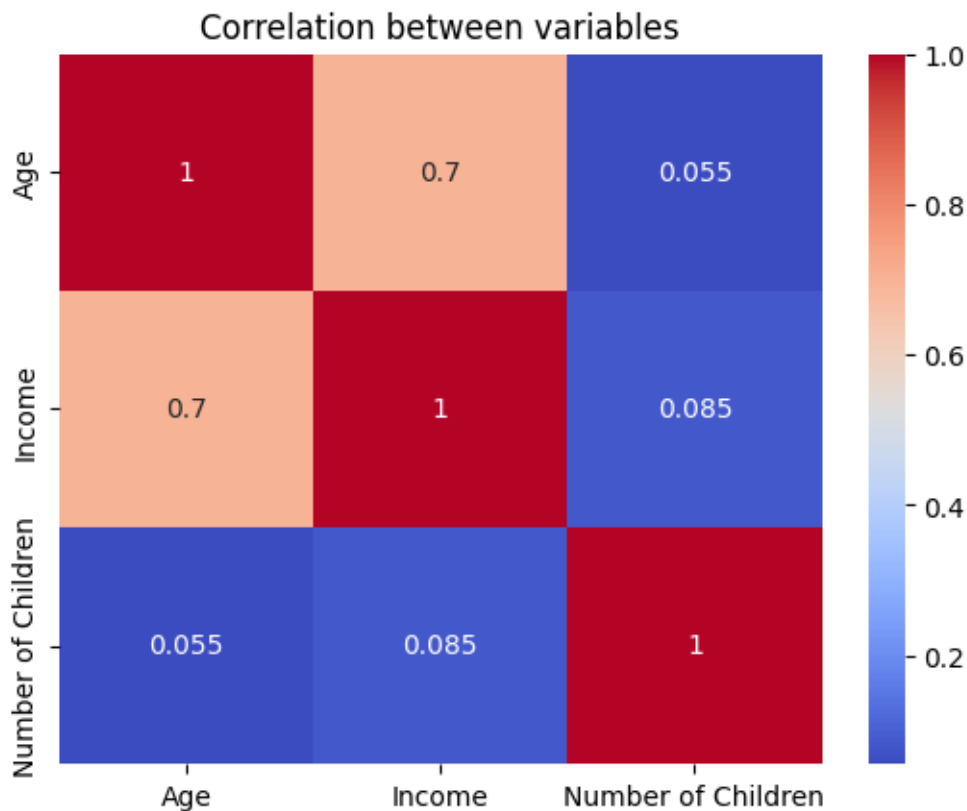
	Age	Income	Number of Children
Age	1.000000	0.699464	0.055390
Income	0.699464	1.000000	0.084547
Number of Children	0.055390	0.084547	1.000000

```
[29]: sns.heatmap(df_original.corr(), annot = True, cmap = 'coolwarm')
plt.title('Correlation between variables')
```

C:\Users\Admin\AppData\Local\Temp\ipykernel_10520\434348305.py:1: FutureWarning:
The default value of numeric_only in DataFrame.corr is deprecated. In a future
version, it will default to False. Select only valid columns or specify the
value of numeric_only to silence this warning.

```
sns.heatmap(df_original.corr(), annot = True, cmap = 'coolwarm')
```


```
[29]: Text(0.5, 1.0, 'Correlation between variables')
```



Heatmap shows that there is a high positive correlation between Age and Income

```
[40]: sns.histplot(df_original, x = df_original['Age'], kde = True)
plt.title('Distribution of Age')
plt.xlabel('Age')
```


```
plt.ylabel('Distribution')
plt.show()
```




```
[47]: df_original['Gender'].value_counts()
```

```
[47]: Female    86
      Male      78
      Name: Gender, dtype: int64
```

```
[52]: sns.barplot(x = df_original['Gender'].unique(), y = df_original['Gender'].
      ↪value_counts())
      plt.title('Value counts of Gender')
      plt.xlabel('Gender')
      plt.ylabel('Count')
      plt.show()
```



```
[54]: sns.histplot(df_original, x = df_original['Income'], kde = True)
plt.title('Distribution of Income')
plt.xlabel('Income')
plt.ylabel('Distribution')
plt.show()
```



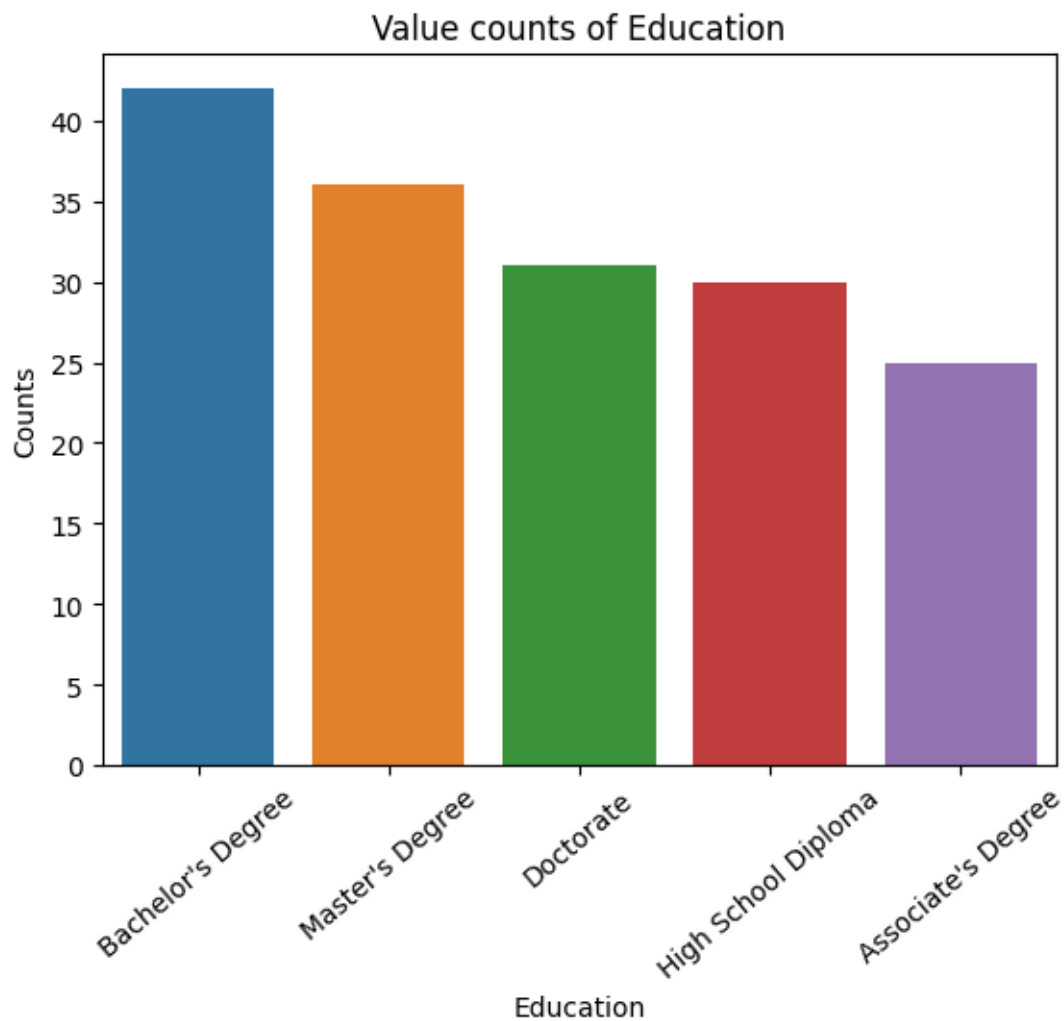
```
[94]: Education = df_original['Education'].value_counts()
Education
```

```
[94]: Bachelor's Degree      42
Master's Degree           36
Doctorate                 31
High School Diploma      30
Associate's Degree        25
Name: Education, dtype: int64
```

```
[101]: sns.barplot(x = Education.index, y = Education)
plt.title('Value counts of Education')
plt.xlabel('Education')
plt.ylabel('Counts')
plt.xticks(rotation = 40)
```

```
[101]: (array([0, 1, 2, 3, 4]),
[Text(0, 0, "Bachelor's Degree"),
Text(1, 0, "Master's Degree"),
Text(2, 0, 'Doctorate'),
```

```
Text(3, 0, 'High School Diploma'),
Text(4, 0, "Associate's Degree"]])
```




```
[86]: Edu_by_Gender = df_original['Education'].groupby(df_original['Gender']).
      ↪value_counts()
      Edu_by_Gender = Edu_by_Gender.reset_index(name = 'Count')
      Edu_by_Gender
```

```
[86]:
```

	Gender	Education	Count
0	Female	Bachelor's Degree	30
1	Female	Doctorate	25
2	Female	High School Diploma	15
3	Female	Associate's Degree	10
4	Female	Master's Degree	6
5	Male	Master's Degree	30

6	Male	Associate's Degree	15
7	Male	High School Diploma	15
8	Male	Bachelor's Degree	12
9	Male	Doctorate	6


```
[104]: sns.barplot(x = Edu_by_Gender['Gender'], y = Edu_by_Gender['Count'], hue = Edu_by_Gender['Education'])
plt.title('Education by Gender')
plt.xlabel('Gender')
plt.ylabel('Counts')
plt.show()
```



```
[110]: df_original['Marital Status'].value_counts()
```

```
[110]: Married    87
Single      77
Name: Marital Status, dtype: int64
```

```
[109]: sns.barplot(x = df_original['Marital Status'].unique(),y = df_original['Marital_
↳Status'].value_counts())
plt.title('Value counts of Marital Status')
plt.xlabel('Marital Status')
plt.ylabel('Counts')
plt.show()
```

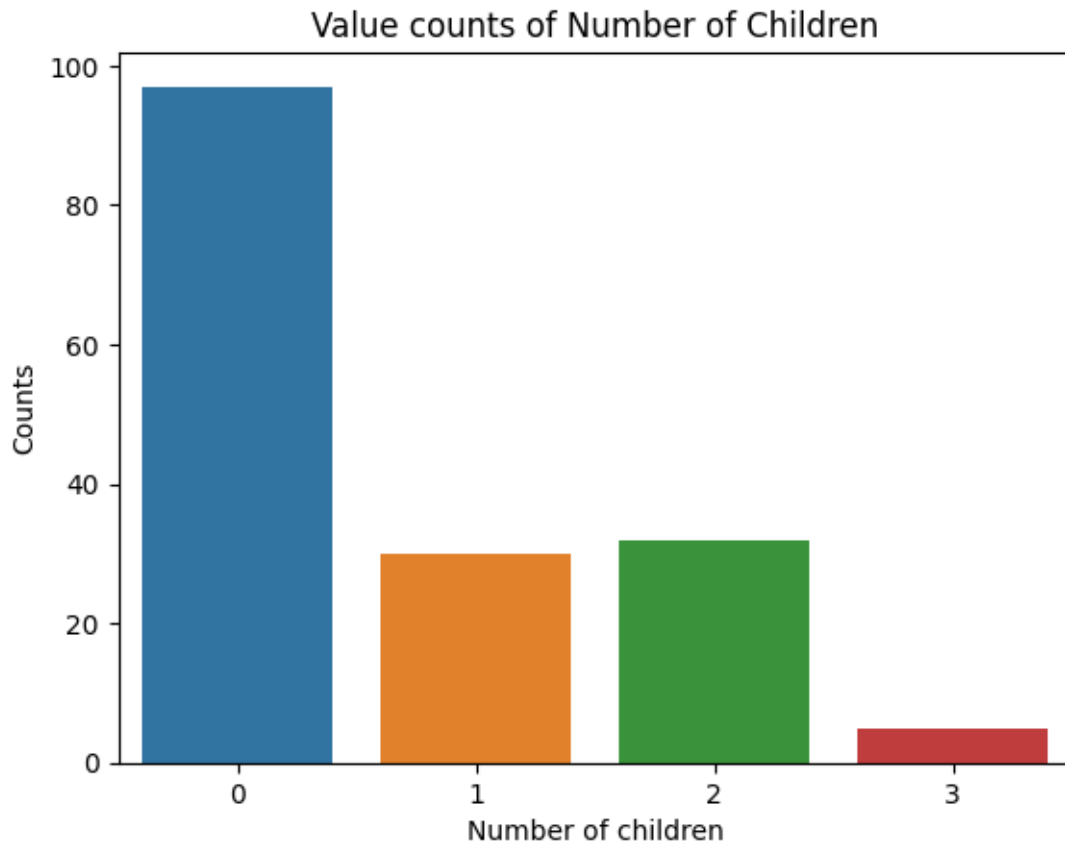


```
[112]: df_original['Number of Children'].value_counts()
```

```
[112]: 0    97
1     32
2     30
3      5
Name: Number of Children, dtype: int64
```

```
[114]: sns.barplot(x = df_original['Number of Children'].unique(), y =
↳df_original['Number of Children'].value_counts())
plt.title('Value counts of Number of Children')
plt.xlabel('Number of children')
plt.ylabel('Counts')
```

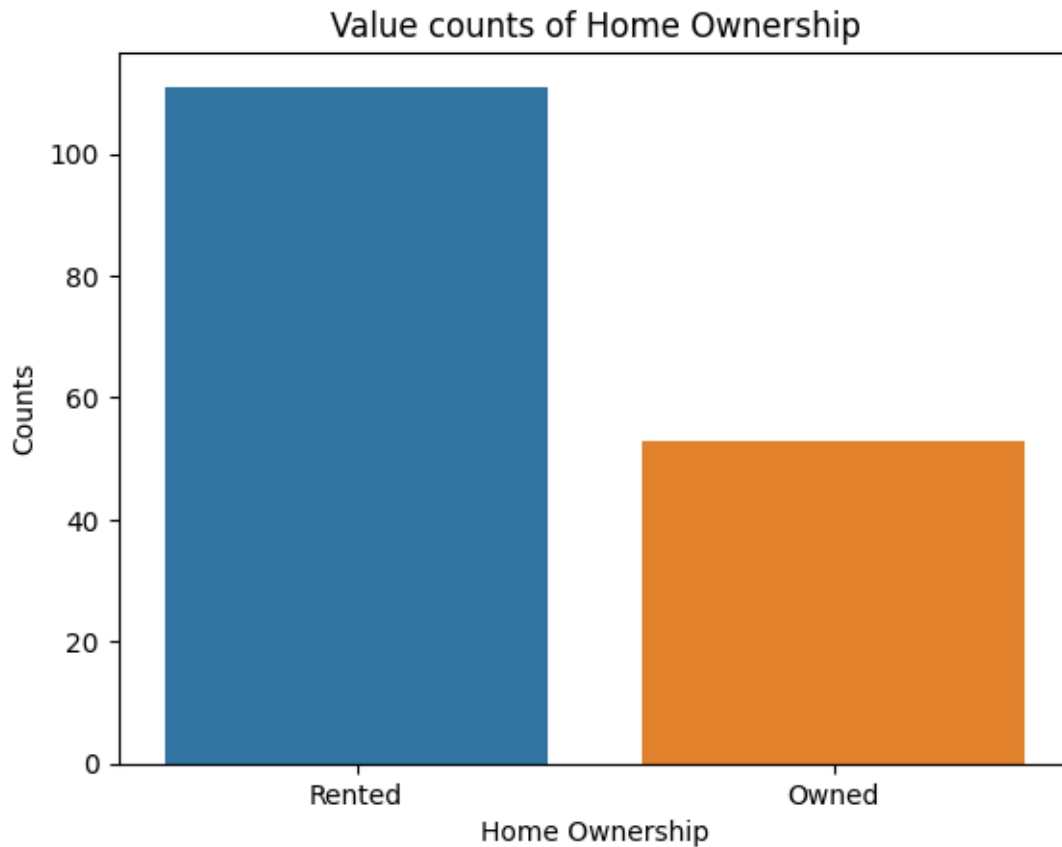
```
plt.show()
```



```
[115]: df_original['Home Ownership'].value_counts()
```

```
[115]: Owned      111  
      Rented     53  
      Name: Home Ownership, dtype: int64
```

```
[117]: sns.barplot(x = df_original['Home Ownership'].unique(), y = df_original['Home_O  
      ↪Ownership'].value_counts())  
      plt.title('Value counts of Home Ownership')  
      plt.xlabel('Home Ownership')  
      plt.ylabel('Counts')  
      plt.show()
```

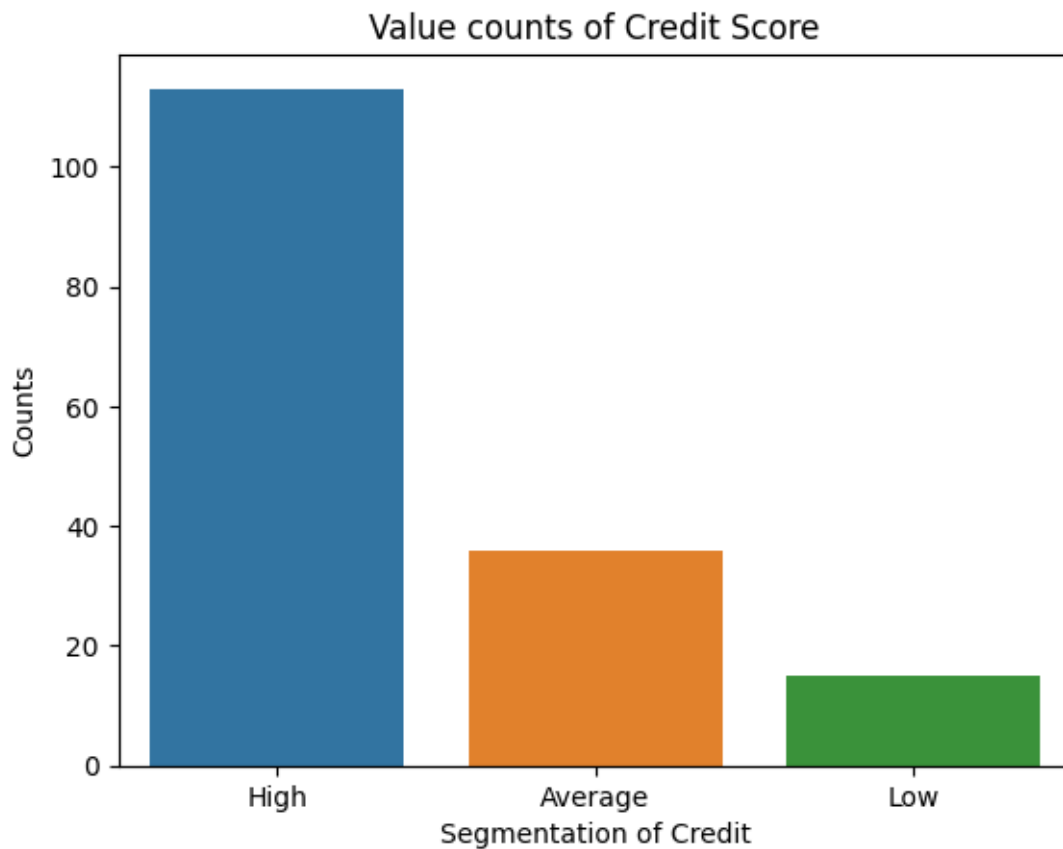


```
[118]: df_original['Credit Score'].value_counts()
```

```
[118]: High      113  
Average   36  
Low       15  
Name: Credit Score, dtype: int64
```

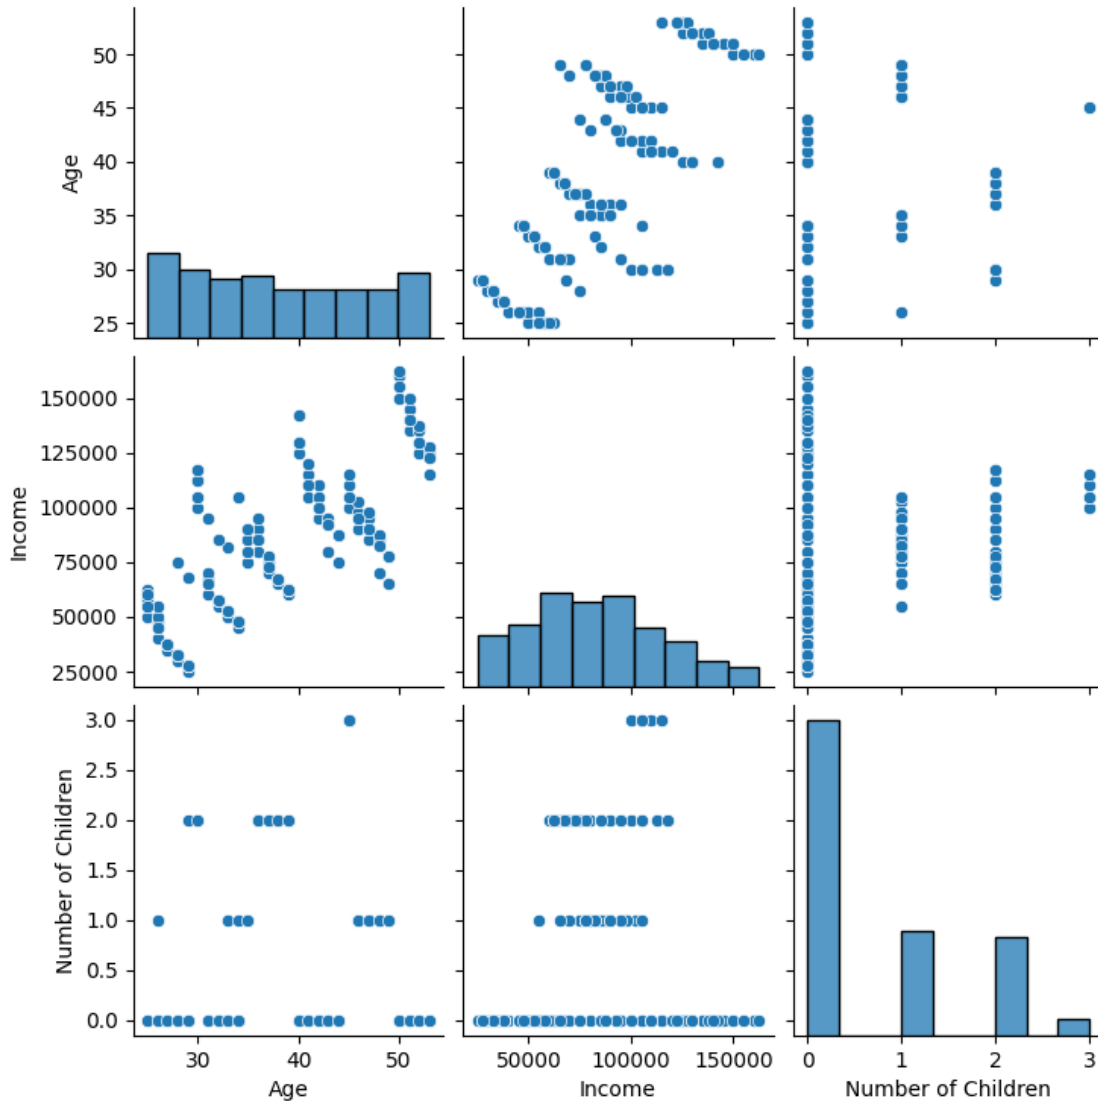
2.4

```
[121]: sns.barplot(x = df_original['Credit Score'].unique(), y = df_original['Credit_  
↪Score'].value_counts())  
plt.title('Value counts of Credit Score')  
plt.xlabel('Segmentation of Credit')  
plt.ylabel('Counts')  
plt.show()
```



```
[123]: sns.pairplot(data = df_original)  
plt.show()
```

```
C:\Users\Admin\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The  
figure layout has changed to tight  
  self._figure.tight_layout(*args, **kwargs)
```



2.5 Label Encoding

```
[4]: le = LabelEncoder()
df_original['Gender'] = le.fit_transform(df_original['Gender'])
df_original['Education'] = le.fit_transform(df_original['Education'])
df_original['Marital Status'] = le.fit_transform(df_original['Marital Status'])
df_original['Home Ownership'] = le.fit_transform(df_original['Home Ownership'])
df_original['Credit Score'] = le.fit_transform(df_original['Credit Score'])
df_original.head()
```

```
[4]:   Age  Gender  Income  Education  Marital Status  Number of Children \
0    25      0   50000          1              1              0
```

1	30	1	100000	4	0	2
2	35	0	75000	2	0	1
3	40	1	125000	3	1	0
4	45	0	100000	1	0	3

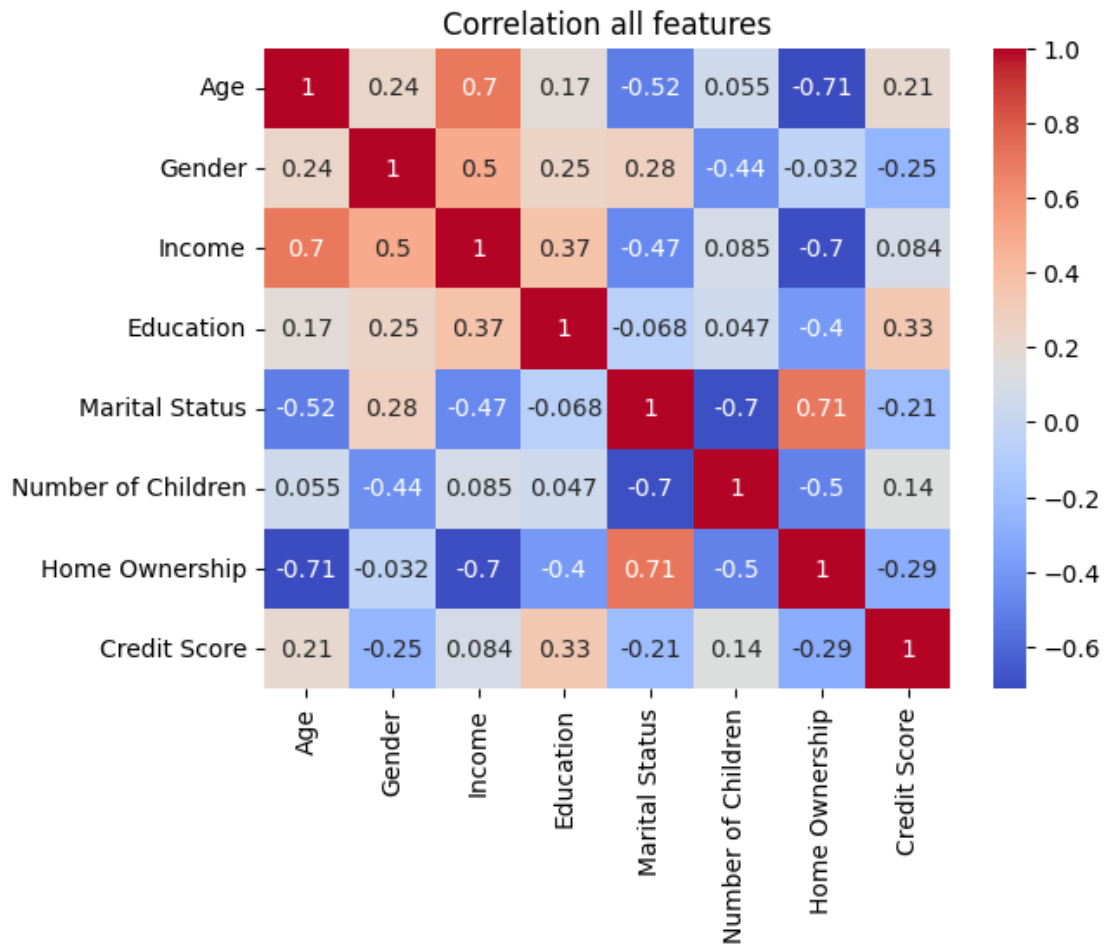
	Home Ownership	Credit Score
0	1	1
1	0	1
2	0	1
3	0	1
4	0	1

I pick out 5 categorical columns to label encoder: Gender, Education, Marital Status, Home Ownership, Credit Score.

2.6 Correlation after encoder

```
[27]: df_original.corr()
sns.heatmap(df_original.corr(), cmap = 'coolwarm', annot = True)
plt.title('Correlation all features')
plt.show
```

```
[27]: <function matplotlib.pyplot.show(close=None, block=None)>
```

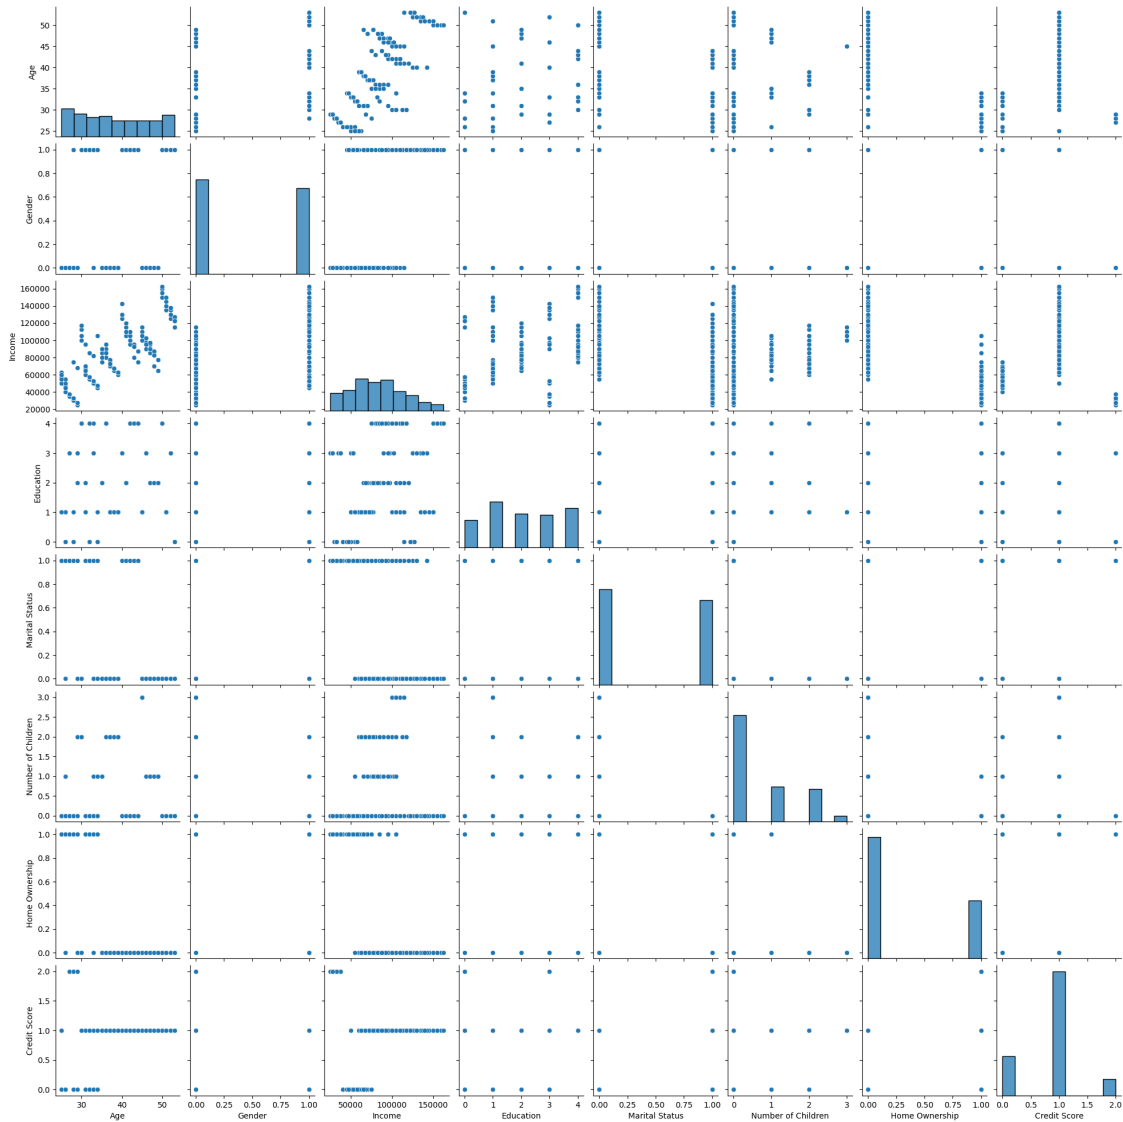


```
[20]: sns.pairplot(df_original)
```

C:\Users\Admin\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight

```
self._figure.tight_layout(*args, **kwargs)
```

```
[20]: <seaborn.axisgrid.PairGrid at 0x19655860dd0>
```

```
[21]: df_original.corr()
```

```
[21]:
```

	Age	Gender	Income	Education	Marital Status	\
Age	1.000000	0.235343	0.699464	0.170254	-0.517723	
Gender	0.235343	1.000000	0.495738	0.248671	0.278362	
Income	0.699464	0.495738	1.000000	0.369449	-0.471004	
Education	0.170254	0.248671	0.369449	1.000000	-0.067797	
Marital Status	-0.517723	0.278362	-0.471004	-0.067797	1.000000	
Number of Children	0.055390	-0.442139	0.084547	0.047311	-0.696984	
Home Ownership	-0.713803	-0.031519	-0.704928	-0.397043	0.708374	
Credit Score	0.205362	-0.247729	0.083698	0.334424	-0.205756	

```
Number of Children Home Ownership Credit Score
```

Age	0.055390	-0.713803	0.205362
Gender	-0.442139	-0.031519	-0.247729
Income	0.084547	-0.704928	0.083698
Education	0.047311	-0.397043	0.334424
Marital Status	-0.696984	0.708374	-0.205756
Number of Children	1.000000	-0.497129	0.136517
Home Ownership	-0.497129	1.000000	-0.293384
Credit Score	0.136517	-0.293384	1.000000

We can see the slightly positive correlation between Education and Credit Score


2.7 Analyzing on target column values 'Credit Score'

```
[32]: df_original['Credit Score'].value_counts()
```

```
[32]: 1    113
      0     36
      2     15
      Name: Credit Score, dtype: int64
```

```
[31]: plt.figure(figsize = (5.5,5.5))
      plt.pie(df_original['Credit Score'].value_counts(), autopct = '%1.2f%%')
      plt.title('Percentage of value in Credit Score')
      plt.xlabel('Credit Score')
      plt.legend(df['Credit Score'].unique())
      plt.show()
```

Percentage of value in Credit Score



Credit Score

There is a significant imbalance between 3 unique values in column Credit Score. I need to address this issue by 3 methods later. I will try 3 methods to find out which is the good method for this dataset

2.7.1 Method 1: Undersampling

```
[115]: df_original['Credit Score'].value_counts()
```

```
[115]: 1    113
      0     36
      2     15
      Name: Credit Score, dtype: int64
```

```
[5]: under_count_0 = df_original[df_original['Credit Score'] == 0]
      under_count_1 = df_original[df_original['Credit Score'] == 1]
      under_count_2 = df_original[df_original['Credit Score'] == 2]
      undersample_size = 15
      undersample_count_0 = under_count_0.sample(undersample_size, random_state = 0)
```

```
undersample_count_1 = under_count_1.sample(undersample_size, random_state = 1)
```

```
[6]: print('Credit 0: ',undersample_count_0.shape)
      print('Credit 1: ',undersample_count_1.shape)
      print('Credit 2: ',under_count_2.shape)
```

```
Credit 0: (15, 8)
Credit 1: (15, 8)
Credit 2: (15, 8)
```

```
[7]: under_df = pd.concat([undersample_count_0, undersample_count_1, under_count_2],
      ↪axis = 0, ignore_index = True)
      under_df.shape
```

```
[7]: (45, 8)
```

```
[7]: under_df.isna().sum()
```

```
[7]: Age                0
      Gender            0
      Income            0
      Education         0
      Marital Status    0
      Number of Children 0
      Home Ownership    0
      Credit Score      0
      dtype: int64
```

```
[125]: sns.barplot(x = under_df['Credit Score'].unique(), y = under_df['Credit Score'].
      ↪value_counts())
      plt.title('Values in Credit Score')
      plt.xlabel('Credit Score Segmentation')
      plt.ylabel('Counts of value')
      plt.show()
```



2.7.2 Method 2: Oversampling

```
[8]: df_original['Credit Score'].value_counts()
```

```
[8]: 1    113  
     0     36  
     2     15  
     Name: Credit Score, dtype: int64
```

```
[9]: over_count_0 = df_original[df_original['Credit Score'] == 0]  
     over_count_1 = df_original[df_original['Credit Score'] == 1]  
     over_count_2 = df_original[df_original['Credit Score'] == 2]  
     oversample_size = 113  
     oversample_count_0 = over_count_0.sample(oversample_size, random_state = 0,   
     ↪replace = True)  
     oversample_count_2 = over_count_2.sample(oversample_size, random_state = 2,   
     ↪replace = True)
```

```
[10]: print('Credit 0: ',oversample_count_0.shape)
      print('Credit 1: ',over_count_1.shape)
      print('Credit 2: ',oversample_count_2.shape)
```

```
Credit 0: (113, 8)
Credit 1: (113, 8)
Credit 2: (113, 8)
```

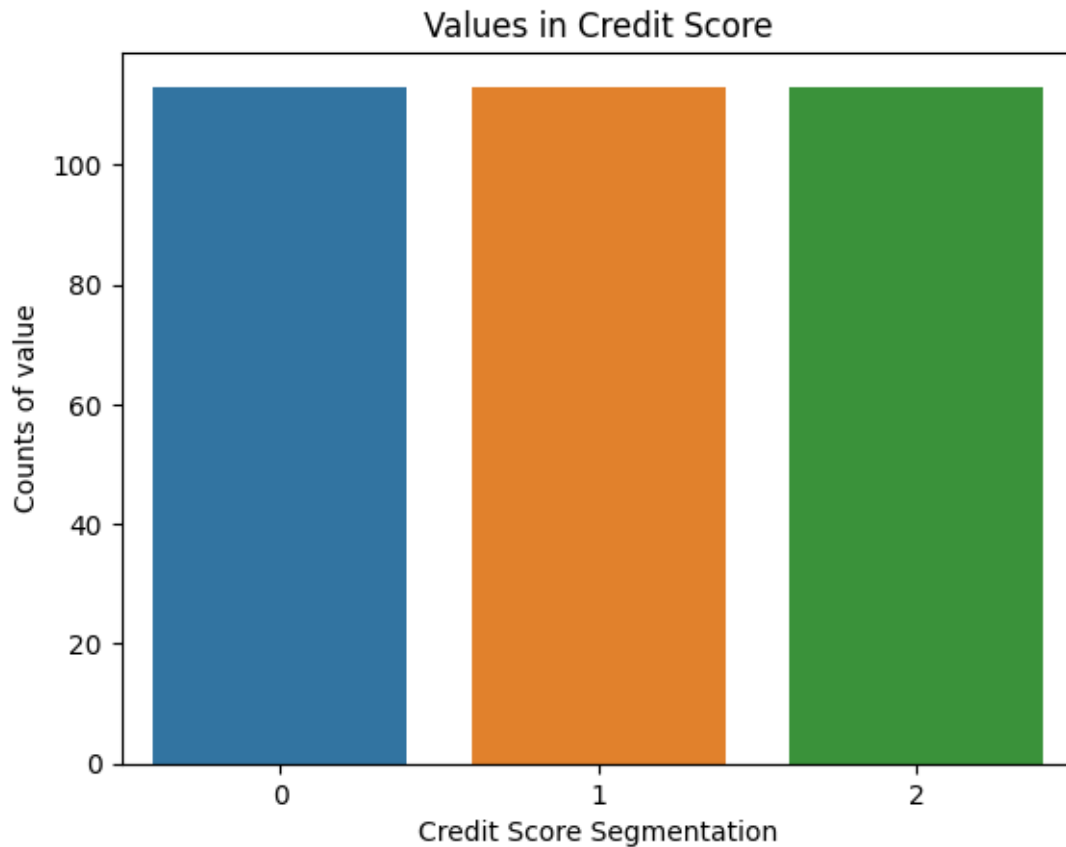
```
[11]: over_df = pd.concat([oversample_count_0, over_count_1, oversample_count_2],
      ↪axis = 0, ignore_index = True)
      over_df.shape
```

```
[11]: (339, 8)
```

```
[12]: over_df.isna().sum()
```

```
[12]: Age                0
      Gender             0
      Income             0
      Education          0
      Marital Status     0
      Number of Children 0
      Home Ownership     0
      Credit Score       0
      dtype: int64
```

```
[139]: sns.barplot(x = over_df['Credit Score'].unique(), y = over_df['Credit Score'].
      ↪value_counts())
      plt.title('Values in Credit Score')
      plt.xlabel('Credit Score Segmentation')
      plt.ylabel('Counts of value')
      plt.show()
```



2.7.3 Method 3: SMOTE

```
[140]: df_original['Credit Score'].value_counts()
```

```
[140]: 1    113
      0     36
      2     15
      Name: Credit Score, dtype: int64
```

```
[13]: !pip install imblearn
      from imblearn.over_sampling import SMOTE
```

[notice] A new release of pip available: 22.3.1 -> 23.2.1

[notice] To update, run: python.exe -m pip install --upgrade pip

Requirement already satisfied: imblearn in c:\users\admin\lib\site-packages (0.0)

Requirement already satisfied: imbalanced-learn in c:\users\admin\lib\site-packages (from imblearn) (0.11.0)

Requirement already satisfied: numpy>=1.17.3 in c:\users\admin\lib\site-packages (from imbalanced-learn->imblearn) (1.24.2)
Requirement already satisfied: scipy>=1.5.0 in c:\users\admin\lib\site-packages (from imbalanced-learn->imblearn) (1.11.1)
Requirement already satisfied: scikit-learn>=1.0.2 in c:\users\admin\lib\site-packages (from imbalanced-learn->imblearn) (1.3.0)
Requirement already satisfied: joblib>=1.1.1 in c:\users\admin\lib\site-packages (from imbalanced-learn->imblearn) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\admin\lib\site-packages (from imbalanced-learn->imblearn) (3.2.0)

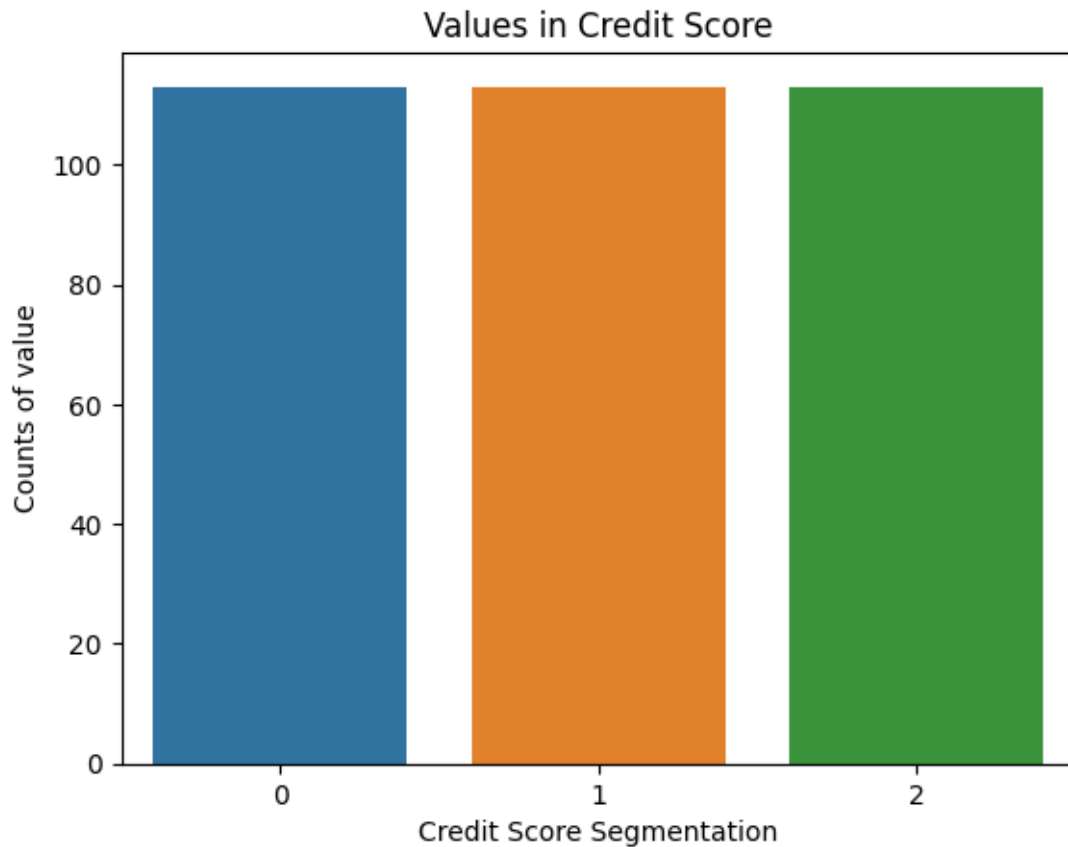
```
[14]: sm = SMOTE(sampling_strategy = 'all', random_state = 3)
X = df_original.drop('Credit Score', axis = 1)
y = df_original['Credit Score']
X_sm, y_sm = sm.fit_resample(X, y)
sm_df = pd.DataFrame(X_sm, columns = df_original.columns)
sm_df['Credit Score'] = y_sm
sm_df.shape
```

```
[14]: (339, 8)
```

```
[13]: sm_df.isna().sum()
```

```
[13]: Age                0
Gender                0
Income                0
Education             0
Marital Status        0
Number of Children    0
Home Ownership         0
Credit Score          0
dtype: int64
```

```
[163]: sns.barplot(x = sm_df['Credit Score'].unique(), y = sm_df['Credit Score'].
↪value_counts())
plt.title('Values in Credit Score')
plt.xlabel('Credit Score Segmentation')
plt.ylabel('Counts of value')
plt.show()
```

2.8 Preprocessing model machine learning

2.8.1 For Method 1: Undersampling

```
[15]: X_under = under_df.drop(['Credit Score'], axis = 1)
      y_under = under_df['Credit Score']
      X_train_under, X_test_under, y_train_under, y_test_under = \
          train_test_split(X_under, y_under, test_size = 0.2, random_state = 10)
```

```
[16]: sc = StandardScaler()
      X_train_under = sc.fit_transform(X_train_under)
      X_test_under = sc.transform(X_test_under)
```

2.8.2 For Method 2: Oversampling

```
[17]: X_over = over_df.drop(['Credit Score'], axis = 1)
      y_over = over_df['Credit Score']
      X_train_over, X_test_over, y_train_over, y_test_over = train_test_split(X_over, \
          y_over, test_size = 0.2, random_state = 10)
```

```
[18]: sc = StandardScaler()
X_train_over = sc.fit_transform(X_train_over)
X_test_over = sc.transform(X_test_over)
```

2.8.3 For Method 3: SMOTE

```
[19]: X_sm = sm_df.drop(['Credit Score'], axis = 1)
y_sm = sm_df['Credit Score']
X_train_sm, X_test_sm, y_train_sm, y_test_sm = train_test_split(X_sm, y_sm,
↳test_size = 0.2, random_state = 10)
```

```
[20]: sc = StandardScaler()
X_train_sm = sc.fit_transform(X_train_sm)
X_test_sm = sc.transform(X_test_sm)
```

2.9 Building model run on 3 methods

```
[27]: lr = LogisticRegression()
dtc = DecisionTreeClassifier()
rfc = RandomForestClassifier()
gbc = GradientBoostingClassifier()
knn = KNeighborsClassifier()
svc = SVC()
gnb = GaussianNB()
xgb = XGBClassifier()
```

```
[29]: models = [lr, dtc, rfc, gbc, knn, svc, gnb, xgb]
```

2.9.1 Method 1: Undersampling

```
[51]: for model in models:
    model.fit(X_train_under, y_train_under)
    y_pred_under = model.predict(X_test_under)
    Classification_rp = classification_report(y_test_under, y_pred_under)
    print('-----')
    print(type(model).__name__)
    print(Classification_rp)
    print(confusion_matrix(y_test_under, y_pred_under))
```

```
-----
LogisticRegression
              precision    recall  f1-score   support

0               0.96         1.00         0.98         27
1               1.00         0.95         0.97         19
2               1.00         1.00         1.00         22
```

accuracy			0.99	68
macro avg	0.99	0.98	0.98	68
weighted avg	0.99	0.99	0.99	68

```
[[2 0 0]
 [0 3 0]
 [0 0 4]]
```

DecisionTreeClassifier

	precision	recall	f1-score	support
0	0.96	1.00	0.98	27
1	1.00	0.95	0.97	19
2	1.00	1.00	1.00	22

accuracy			0.99	68
macro avg	0.99	0.98	0.98	68
weighted avg	0.99	0.99	0.99	68

```
[[2 0 0]
 [1 2 0]
 [0 0 4]]
```

RandomForestClassifier

	precision	recall	f1-score	support
0	0.96	1.00	0.98	27
1	1.00	0.95	0.97	19
2	1.00	1.00	1.00	22

accuracy			0.99	68
macro avg	0.99	0.98	0.98	68
weighted avg	0.99	0.99	0.99	68

```
[[2 0 0]
 [0 3 0]
 [0 0 4]]
```

GradientBoostingClassifier

	precision	recall	f1-score	support
0	0.96	1.00	0.98	27
1	1.00	0.95	0.97	19
2	1.00	1.00	1.00	22

accuracy			0.99	68
macro avg	0.99	0.98	0.98	68

weighted avg	0.99	0.99	0.99	68
--------------	------	------	------	----

```
[[2 0 0]
 [1 2 0]
 [0 0 4]]
```

KNeighborsClassifier

	precision	recall	f1-score	support
0	0.96	1.00	0.98	27
1	1.00	0.95	0.97	19
2	1.00	1.00	1.00	22
accuracy			0.99	68
macro avg	0.99	0.98	0.98	68
weighted avg	0.99	0.99	0.99	68

```
[[2 0 0]
 [0 3 0]
 [0 0 4]]
```

SVC

	precision	recall	f1-score	support
0	0.96	1.00	0.98	27
1	1.00	0.95	0.97	19
2	1.00	1.00	1.00	22
accuracy			0.99	68
macro avg	0.99	0.98	0.98	68
weighted avg	0.99	0.99	0.99	68

```
[[1 0 1]
 [0 3 0]
 [0 0 4]]
```

GaussianNB

	precision	recall	f1-score	support
0	0.96	1.00	0.98	27
1	1.00	0.95	0.97	19
2	1.00	1.00	1.00	22
accuracy			0.99	68
macro avg	0.99	0.98	0.98	68
weighted avg	0.99	0.99	0.99	68

```
[[1 0 1]
```

```
[1 2 0]
[0 0 4]]
```

XGBClassifier

	precision	recall	f1-score	support
0	0.96	1.00	0.98	27
1	1.00	0.95	0.97	19
2	1.00	1.00	1.00	22
accuracy			0.99	68
macro avg	0.99	0.98	0.98	68
weighted avg	0.99	0.99	0.99	68

```
[[2 0 0]
 [0 3 0]
 [0 0 4]]
```

For this undersampling method, we see all models give good outcomes.

2.9.2 Method 2: Oversampling

```
[52]: for model in models:
        model.fit(X_train_over, y_train_over)
        y_pred_over = model.predict(X_test_over)
        Classification_rp = classification_report(y_test_over, y_pred_over)
        print('-----')
        print(type(model).__name__)
        print(Classification_rp)
        print(confusion_matrix(y_test_over, y_pred_over))
```

LogisticRegression

	precision	recall	f1-score	support
0	0.86	0.95	0.90	19
1	0.96	0.89	0.93	28
2	1.00	1.00	1.00	21
accuracy			0.94	68
macro avg	0.94	0.95	0.94	68
weighted avg	0.94	0.94	0.94	68

```
[[18 1 0]
 [ 3 25 0]
 [ 0 0 21]]
```

DecisionTreeClassifier

	precision	recall	f1-score	support
0	0.90	1.00	0.95	19
1	1.00	0.93	0.96	28
2	1.00	1.00	1.00	21
accuracy			0.97	68
macro avg	0.97	0.98	0.97	68
weighted avg	0.97	0.97	0.97	68

[[19 0 0]
[2 26 0]
[0 0 21]]

RandomForestClassifier

	precision	recall	f1-score	support
0	0.90	1.00	0.95	19
1	1.00	0.93	0.96	28
2	1.00	1.00	1.00	21
accuracy			0.97	68
macro avg	0.97	0.98	0.97	68
weighted avg	0.97	0.97	0.97	68

[[19 0 0]
[2 26 0]
[0 0 21]]

GradientBoostingClassifier

	precision	recall	f1-score	support
0	0.86	1.00	0.93	19
1	1.00	0.89	0.94	28
2	1.00	1.00	1.00	21
accuracy			0.96	68
macro avg	0.95	0.96	0.96	68
weighted avg	0.96	0.96	0.96	68

[[19 0 0]
[3 25 0]
[0 0 21]]

KNeighborsClassifier

	precision	recall	f1-score	support
0	0.86	0.95	0.90	19

1	0.96	0.89	0.93	28
2	1.00	1.00	1.00	21
accuracy			0.94	68
macro avg	0.94	0.95	0.94	68
weighted avg	0.94	0.94	0.94	68

```
[[18  1  0]
 [ 3 25  0]
 [ 0  0 21]]
```

SVC

	precision	recall	f1-score	support
0	0.89	0.84	0.86	19
1	0.96	0.93	0.95	28
2	0.91	1.00	0.95	21
accuracy			0.93	68
macro avg	0.92	0.92	0.92	68
weighted avg	0.93	0.93	0.93	68

```
[[16  1  2]
 [ 2 26  0]
 [ 0  0 21]]
```

GaussianNB

	precision	recall	f1-score	support
0	0.92	0.58	0.71	19
1	0.96	0.93	0.95	28
2	0.72	1.00	0.84	21
accuracy			0.85	68
macro avg	0.87	0.84	0.83	68
weighted avg	0.88	0.85	0.85	68

```
[[11  1  7]
 [ 1 26  1]
 [ 0  0 21]]
```

XGBClassifier

	precision	recall	f1-score	support
0	0.95	1.00	0.97	19
1	1.00	0.96	0.98	28
2	1.00	1.00	1.00	21

accuracy			0.99	68
macro avg	0.98	0.99	0.99	68
weighted avg	0.99	0.99	0.99	68

```
[[19  0  0]
 [ 1 27  0]
 [ 0  0 21]]
```

With Oversampling method, although models can give out good outcomes, we clearly see that the outcomes are not as good as in comparison with Undersampling method. In this sample method, XG Boost Classifier is the best performance model.

2.9.3 Method 3: SMOTE

```
[53]: for model in models:
        model.fit(X_train_sm, y_train_sm)
        y_pred_sm = model.predict(X_test_sm)
        Classification_rp = classification_report(y_test_sm, y_pred_sm)
        print('-----')
        print(type(model).__name__)
        print(Classification_rp)
        print(confusion_matrix(y_test_sm, y_pred_sm))
```

```
-----
LogisticRegression
      precision    recall  f1-score   support

    0       0.96      0.96      0.96         27
    1       1.00      0.95      0.97         19
    2       0.96      1.00      0.98         22

 accuracy
macro avg       0.97      0.97      0.97         68
weighted avg       0.97      0.97      0.97         68

[[26  0  1]
 [ 1 18  0]
 [ 0  0 22]]
```

```
-----
DecisionTreeClassifier
      precision    recall  f1-score   support

    0       0.96      1.00      0.98         27
    1       1.00      0.95      0.97         19
    2       1.00      1.00      1.00         22

 accuracy
macro avg       0.99      0.98      0.98         68
```


weighted avg	0.99	0.99	0.99	68
--------------	------	------	------	----

```
[[27  0  0]
 [ 1 18  0]
 [ 0  0 22]]
```

RandomForestClassifier

	precision	recall	f1-score	support
0	0.96	1.00	0.98	27
1	1.00	0.95	0.97	19
2	1.00	1.00	1.00	22
accuracy			0.99	68
macro avg	0.99	0.98	0.98	68
weighted avg	0.99	0.99	0.99	68

```
[[27  0  0]
 [ 1 18  0]
 [ 0  0 22]]
```

GradientBoostingClassifier

	precision	recall	f1-score	support
0	1.00	0.93	0.96	27
1	0.90	1.00	0.95	19
2	1.00	1.00	1.00	22
accuracy			0.97	68
macro avg	0.97	0.98	0.97	68
weighted avg	0.97	0.97	0.97	68

```
[[25  2  0]
 [ 0 19  0]
 [ 0  0 22]]
```

KNeighborsClassifier

	precision	recall	f1-score	support
0	0.96	1.00	0.98	27
1	1.00	0.95	0.97	19
2	1.00	1.00	1.00	22
accuracy			0.99	68
macro avg	0.99	0.98	0.98	68
weighted avg	0.99	0.99	0.99	68

```
[[27  0  0]
```

```
[ 1 18  0]
[ 0  0 22]]
```

SVC

	precision	recall	f1-score	support
0	0.96	0.93	0.94	27
1	1.00	0.95	0.97	19
2	0.92	1.00	0.96	22
accuracy			0.96	68
macro avg	0.96	0.96	0.96	68
weighted avg	0.96	0.96	0.96	68

```
[[25  0  2]
 [ 1 18  0]
 [ 0  0 22]]
```

GaussianNB

	precision	recall	f1-score	support
0	0.93	0.52	0.67	27
1	1.00	0.95	0.97	19
2	0.63	1.00	0.77	22
accuracy			0.79	68
macro avg	0.85	0.82	0.80	68
weighted avg	0.85	0.79	0.79	68

```
[[14  0 13]
 [ 1 18  0]
 [ 0  0 22]]
```

XGBClassifier

	precision	recall	f1-score	support
0	1.00	1.00	1.00	27
1	1.00	1.00	1.00	19
2	1.00	1.00	1.00	22
accuracy			1.00	68
macro avg	1.00	1.00	1.00	68
weighted avg	1.00	1.00	1.00	68

```
[[27  0  0]
 [ 0 19  0]
 [ 0  0 22]]
```

In SMOTE method, all models give very good results, possibly on par with method Undersampling. Particularly, the XGB model across the three methods consistently gives the best outcomes among the models, and in this method, XGB achieves the best performance in all three runs.

2.10 In Conclusion

This is really an interesting project that I have done. After analyzing, I give the following two conclusions: - About the method of handling imbalanced data: I recommend choosing the SMOTE method - About the machine learning model: I recommend choosing the XG Boost Classifier model

Thank you for reading my personal project

Author: Bao Thai