

fraud-credit-card-detection

August 21, 2023

1 About Dataset

1.1 Context

It is important that credit card companies are able to recognize fraudulent credit card transactions so that customers are not charged for items that they did not purchase.

1.2 Content

The dataset contains transactions made by credit cards in September 2013 by European cardholders. This dataset presents transactions that occurred in two days, where we have 492 frauds out of 284,807 transactions. The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions.

It contains only numerical input variables which are the result of a PCA transformation. Unfortunately, due to confidentiality issues, we cannot provide the original features and more background information about the data. Features V1, V2, ... V28 are the principal components obtained with PCA, the only features which have not been transformed with PCA are 'Time' and 'Amount'. Feature 'Time' contains the seconds elapsed between each transaction and the first transaction in the dataset. The feature 'Amount' is the transaction Amount, this feature can be used for example-dependant cost-sensitive learning. Feature 'Class' is the response variable and it takes value 1 in case of fraud and 0 otherwise.

Given the class imbalance ratio, we recommend measuring the accuracy using the Area Under the Precision-Recall Curve (AUPRC). Confusion matrix accuracy is not meaningful for unbalanced classification.

1.3 Source and link

Source: Kaggle Link: <https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud>

Author: Bao Thai

2 Import needed libraries and packages

```
[1]: import numpy as np
import pandas as pd
!pip install matplotlib
import matplotlib.pyplot as plt
!pip install seaborn
```

```

import seaborn as sns
!pip install scipy
from scipy import stats
!pip install scikit-learn
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_val_score
from sklearn.metrics import confusion_matrix, accuracy_score,
    classification_report
from sklearn.metrics import r2_score, f1_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
!pip install xgboost
from xgboost import XGBClassifier
from csv import reader

```

[notice] A new release of pip available: 22.3.1 -> 23.2.1

[notice] To update, run: python.exe -m pip install --upgrade pip

Requirement already satisfied: matplotlib in c:\users\admin\lib\site-packages (3.7.2)

Requirement already satisfied: contourpy>=1.0.1 in c:\users\admin\lib\site-packages (from matplotlib) (1.0.7)

Requirement already satisfied: cycler>=0.10 in c:\users\admin\lib\site-packages (from matplotlib) (0.11.0)

Requirement already satisfied: fonttools>=4.22.0 in c:\users\admin\lib\site-packages (from matplotlib) (4.39.3)

Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\admin\lib\site-packages (from matplotlib) (1.4.4)

Requirement already satisfied: numpy>=1.20 in c:\users\admin\lib\site-packages (from matplotlib) (1.24.2)

Requirement already satisfied: packaging>=20.0 in c:\users\admin\lib\site-packages (from matplotlib) (23.0)

Requirement already satisfied: pillow>=6.2.0 in c:\users\admin\lib\site-packages (from matplotlib) (9.5.0)

Requirement already satisfied: pyparsing<3.1,>=2.3.1 in c:\users\admin\lib\site-packages (from matplotlib) (3.0.9)

Requirement already satisfied: python-dateutil>=2.7 in c:\users\admin\lib\site-packages (from matplotlib) (2.8.2)

Requirement already satisfied: six>=1.5 in c:\users\admin\lib\site-packages (from python-dateutil>=2.7->matplotlib) (1.16.0)

[notice] A new release of pip available: 22.3.1 -> 23.2.1

[notice] To update, run: python.exe -m pip install --upgrade pip

Requirement already satisfied: seaborn in c:\users\admin\lib\site-packages (0.12.2)

Requirement already satisfied: numpy!=1.24.0,>=1.17 in c:\users\admin\lib\site-packages (from seaborn) (1.24.2)

Requirement already satisfied: pandas>=0.25 in c:\users\admin\lib\site-packages (from seaborn) (1.5.3)

Requirement already satisfied: matplotlib!=3.6.1,>=3.1 in c:\users\admin\lib\site-packages (from seaborn) (3.7.2)

Requirement already satisfied: contourpy>=1.0.1 in c:\users\admin\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.0.7)

Requirement already satisfied: cycler>=0.10 in c:\users\admin\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (0.11.0)

Requirement already satisfied: fonttools>=4.22.0 in c:\users\admin\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (4.39.3)

Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\admin\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.4.4)

Requirement already satisfied: packaging>=20.0 in c:\users\admin\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (23.0)

Requirement already satisfied: pillow>=6.2.0 in c:\users\admin\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (9.5.0)

Requirement already satisfied: pyparsing<3.1,>=2.3.1 in c:\users\admin\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (3.0.9)

Requirement already satisfied: python-dateutil>=2.7 in c:\users\admin\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in c:\users\admin\lib\site-packages (from pandas>=0.25->seaborn) (2022.7.1)

Requirement already satisfied: six>=1.5 in c:\users\admin\lib\site-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.1->seaborn) (1.16.0)

Requirement already satisfied: scipy in c:\users\admin\lib\site-packages (1.11.1)

[notice] A new release of pip available: 22.3.1 -> 23.2.1

[notice] To update, run: python.exe -m pip install --upgrade pip

Requirement already satisfied: numpy<1.28.0,>=1.21.6 in c:\users\admin\lib\site-packages (from scipy) (1.24.2)

Requirement already satisfied: scikit-learn in c:\users\admin\lib\site-packages (1.3.0)

Requirement already satisfied: numpy>=1.17.3 in c:\users\admin\lib\site-packages (from scikit-learn) (1.24.2)

Requirement already satisfied: scipy>=1.5.0 in c:\users\admin\lib\site-packages (from scikit-learn) (1.11.1)

Requirement already satisfied: joblib>=1.1.1 in c:\users\admin\lib\site-packages

```
(from scikit-learn) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\admin\lib\site-packages (from scikit-learn) (3.2.0)
```

```
[notice] A new release of pip available: 22.3.1 -> 23.2.1
[notice] To update, run: python.exe -m pip install --upgrade pip
```

```
Requirement already satisfied: xgboost in c:\users\admin\lib\site-packages (1.7.6)
```

```
Requirement already satisfied: numpy in c:\users\admin\lib\site-packages (from xgboost) (1.24.2)
```

```
Requirement already satisfied: scipy in c:\users\admin\lib\site-packages (from xgboost) (1.11.1)
```

```
[notice] A new release of pip available: 22.3.1 -> 23.2.1
[notice] To update, run: python.exe -m pip install --upgrade pip
```

2.1 Loading Dataset

```
[2]: data_original = pd.read_csv('C:/Users/Admin/Desktop/Projects/creditcard.csv')
df = data_original.copy()
df.head(20)
```

```
[2]:
```

	Time	V1	V2	V3	V4	V5	V6	\
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	
5	2.0	-0.425966	0.960523	1.141109	-0.168252	0.420987	-0.029728	
6	4.0	1.229658	0.141004	0.045371	1.202613	0.191881	0.272708	
7	7.0	-0.644269	1.417964	1.074380	-0.492199	0.948934	0.428118	
8	7.0	-0.894286	0.286157	-0.113192	-0.271526	2.669599	3.721818	
9	9.0	-0.338262	1.119593	1.044367	-0.222187	0.499361	-0.246761	
10	10.0	1.449044	-1.176339	0.913860	-1.375667	-1.971383	-0.629152	
11	10.0	0.384978	0.616109	-0.874300	-0.094019	2.924584	3.317027	
12	10.0	1.249999	-1.221637	0.383930	-1.234899	-1.485419	-0.753230	
13	11.0	1.069374	0.287722	0.828613	2.712520	-0.178398	0.337544	
14	12.0	-2.791855	-0.327771	1.641750	1.767473	-0.136588	0.807596	
15	12.0	-0.752417	0.345485	2.057323	-1.468643	-1.158394	-0.077850	
16	12.0	1.103215	-0.040296	1.267332	1.289091	-0.735997	0.288069	
17	13.0	-0.436905	0.918966	0.924591	-0.727219	0.915679	-0.127867	
18	14.0	-5.401258	-5.450148	1.186305	1.736239	3.049106	-1.763406	
19	15.0	1.492936	-1.029346	0.454795	-1.438026	-1.555434	-0.720961	

	V7	V8	V9	...	V21	V22	V23	V24	\
0	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	

1	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846
2	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281
3	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575
4	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267
5	0.476201	0.260314	-0.568671	...	-0.208254	-0.559825	-0.026398	-0.371427
6	-0.005159	0.081213	0.464960	...	-0.167716	-0.270710	-0.154104	-0.780055
7	1.120631	-3.807864	0.615375	...	1.943465	-1.015455	0.057504	-0.649709
8	0.370145	0.851084	-0.392048	...	-0.073425	-0.268092	-0.204233	1.011592
9	0.651583	0.069539	-0.736727	...	-0.246914	-0.633753	-0.120794	-0.385050
10	-1.423236	0.048456	-1.720408	...	-0.009302	0.313894	0.027740	0.500512
11	0.470455	0.538247	-0.558895	...	0.049924	0.238422	0.009130	0.996710
12	-0.689405	-0.227487	-2.094011	...	-0.231809	-0.483285	0.084668	0.392831
13	-0.096717	0.115982	-0.221083	...	-0.036876	0.074412	-0.071407	0.104744
14	-0.422911	-1.907107	0.755713	...	1.151663	0.222182	1.020586	0.028317
15	-0.608581	0.003603	-0.436167	...	0.499625	1.353650	-0.256573	-0.065084
16	-0.586057	0.189380	0.782333	...	-0.024612	0.196002	0.013802	0.103758
17	0.707642	0.087962	-0.665271	...	-0.194796	-0.672638	-0.156858	-0.888386
18	-1.559738	0.160842	1.233090	...	-0.503600	0.984460	2.458589	0.042119
19	-1.080664	-0.053127	-1.978682	...	-0.177650	-0.175074	0.040002	0.295814

	V25	V26	V27	V28	Amount	Class
0	0.128539	-0.189115	0.133558	-0.021053	149.62	0
1	0.167170	0.125895	-0.008983	0.014724	2.69	0
2	-0.327642	-0.139097	-0.055353	-0.059752	378.66	0
3	0.647376	-0.221929	0.062723	0.061458	123.50	0
4	-0.206010	0.502292	0.219422	0.215153	69.99	0
5	-0.232794	0.105915	0.253844	0.081080	3.67	0
6	0.750137	-0.257237	0.034507	0.005168	4.99	0
7	-0.415267	-0.051634	-1.206921	-1.085339	40.80	0
8	0.373205	-0.384157	0.011747	0.142404	93.20	0
9	-0.069733	0.094199	0.246219	0.083076	3.68	0
10	0.251367	-0.129478	0.042850	0.016253	7.80	0
11	-0.767315	-0.492208	0.042472	-0.054337	9.99	0
12	0.161135	-0.354990	0.026416	0.042422	121.50	0
13	0.548265	0.104094	0.021491	0.021293	27.50	0
14	-0.232746	-0.235557	-0.164778	-0.030154	58.80	0
15	-0.039124	-0.087086	-0.180998	0.129394	15.99	0
16	0.364298	-0.382261	0.092809	0.037051	12.99	0
17	-0.342413	-0.049027	0.079692	0.131024	0.89	0
18	-0.481631	-0.621272	0.392053	0.949594	46.80	0
19	0.332931	-0.220385	0.022298	0.007602	5.00	0

[20 rows x 31 columns]

2.2 Explore Data

```
[53]: df.shape
```

```
[53]: (284807, 31)
```

```
[54]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0   Time    284807 non-null  float64
 1   V1      284807 non-null  float64
 2   V2      284807 non-null  float64
 3   V3      284807 non-null  float64
 4   V4      284807 non-null  float64
 5   V5      284807 non-null  float64
 6   V6      284807 non-null  float64
 7   V7      284807 non-null  float64
 8   V8      284807 non-null  float64
 9   V9      284807 non-null  float64
10  V10     284807 non-null  float64
11  V11     284807 non-null  float64
12  V12     284807 non-null  float64
13  V13     284807 non-null  float64
14  V14     284807 non-null  float64
15  V15     284807 non-null  float64
16  V16     284807 non-null  float64
17  V17     284807 non-null  float64
18  V18     284807 non-null  float64
19  V19     284807 non-null  float64
20  V20     284807 non-null  float64
21  V21     284807 non-null  float64
22  V22     284807 non-null  float64
23  V23     284807 non-null  float64
24  V24     284807 non-null  float64
25  V25     284807 non-null  float64
26  V26     284807 non-null  float64
27  V27     284807 non-null  float64
28  V28     284807 non-null  float64
29  Amount  284807 non-null  float64
30  Class   284807 non-null  int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

```
[55]: df.describe()
```

[55]:

	Time	V1	V2	V3	V4 \
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15	2.074095e-15
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01

	V5	V6	V7	V8	V9 \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	9.604066e-16	1.487313e-15	-5.556467e-16	1.213481e-16	-2.406331e-15
std	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00	1.098632e+00
min	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01	-1.343407e+01
25%	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086297e-01	-6.430976e-01
50%	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235804e-02	-5.142873e-02
75%	6.119264e-01	3.985649e-01	5.704361e-01	3.273459e-01	5.971390e-01
max	3.480167e+01	7.330163e+01	1.205895e+02	2.000721e+01	1.559499e+01

	...	V21	V22	V23	V24 \
count	...	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	...	1.654067e-16	-3.568593e-16	2.578648e-16	4.473266e-15
std	...	7.345240e-01	7.257016e-01	6.244603e-01	6.056471e-01
min	...	-3.483038e+01	-1.093314e+01	-4.480774e+01	-2.836627e+00
25%	...	-2.283949e-01	-5.423504e-01	-1.618463e-01	-3.545861e-01
50%	...	-2.945017e-02	6.781943e-03	-1.119293e-02	4.097606e-02
75%	...	1.863772e-01	5.285536e-01	1.476421e-01	4.395266e-01
max	...	2.720284e+01	1.050309e+01	2.252841e+01	4.584549e+00

	V25	V26	V27	V28	Amount \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	284807.000000
mean	5.340915e-16	1.683437e-15	-3.660091e-16	-1.227390e-16	88.349619
std	5.212781e-01	4.822270e-01	4.036325e-01	3.300833e-01	250.120109
min	-1.029540e+01	-2.604551e+00	-2.256568e+01	-1.543008e+01	0.000000
25%	-3.171451e-01	-3.269839e-01	-7.083953e-02	-5.295979e-02	5.600000
50%	1.659350e-02	-5.213911e-02	1.342146e-03	1.124383e-02	22.000000
75%	3.507156e-01	2.409522e-01	9.104512e-02	7.827995e-02	77.165000
max	7.519589e+00	3.517346e+00	3.161220e+01	3.384781e+01	25691.160000

	Class
count	284807.000000
mean	0.001727
std	0.041527
min	0.000000
25%	0.000000
50%	0.000000

```
75%          0.000000
max          1.000000
```

```
[8 rows x 31 columns]
```

```
[56]: df.isna().sum()
```

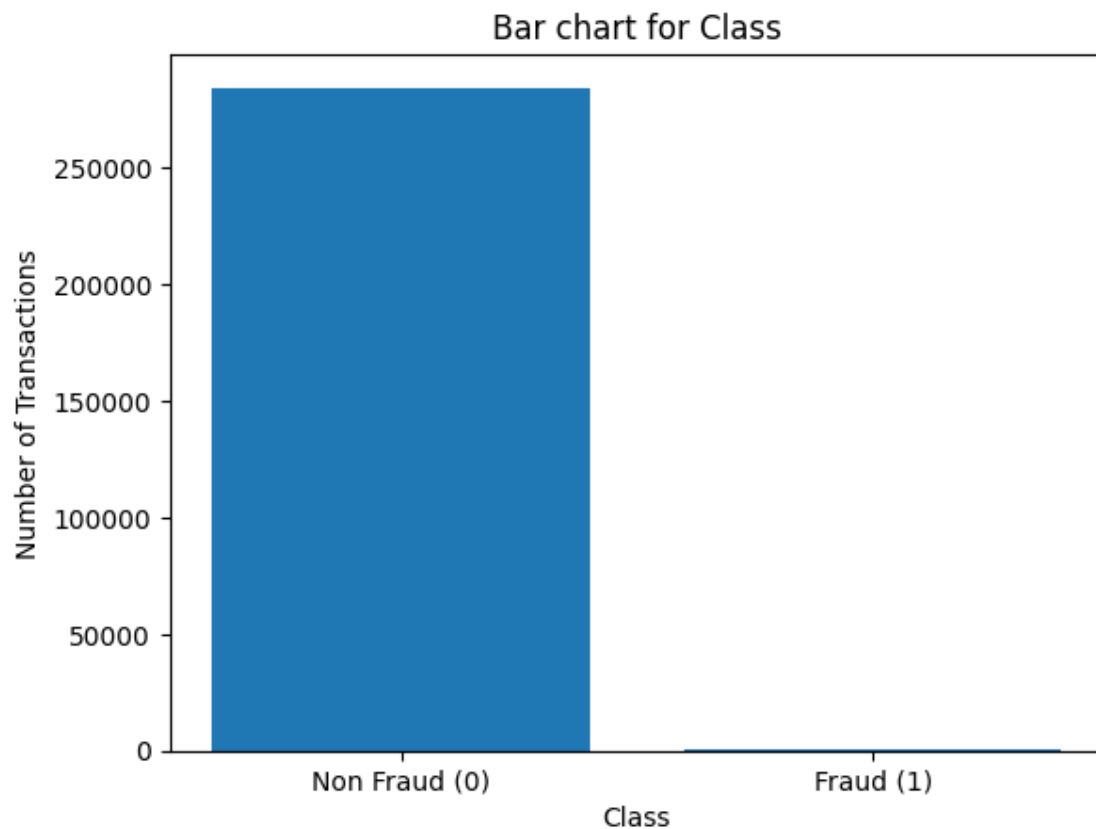
```
[56]: Time      0
      V1        0
      V2        0
      V3        0
      V4        0
      V5        0
      V6        0
      V7        0
      V8        0
      V9        0
      V10       0
      V11       0
      V12       0
      V13       0
      V14       0
      V15       0
      V16       0
      V17       0
      V18       0
      V19       0
      V20       0
      V21       0
      V22       0
      V23       0
      V24       0
      V25       0
      V26       0
      V27       0
      V28       0
      Amount    0
      Class     0
      dtype: int64
```

```
[4]: Class_count = df['Class'].value_counts()
      Class_count
```

```
[4]: 0    284315
      1     492
      Name: Class, dtype: int64
```



```
[5]: plt.bar(Class_count.index, Class_count)
plt.title('Bar chart for Class')
plt.xlabel('Class')
plt.ylabel('Number of Transactions')
plt.xticks([0,1],['Non Fraud (0)', 'Fraud (1)'])
plt.show()
```



This bar chart illustrates the target column named ‘Class’ with a significant imbalance between two values: 0 and 1, Non Fraud and Fraud respectively. I will address this issue later

2.3 Preprocessing and Visualizing

```
[3]: df['Time'] = pd.to_datetime(df['Time'], unit = 'h')
df['Time'] = df['Time'].dt.hour
df.tail()
```

```
[3]:
```

	Time	V1	V2	V3	V4	V5	V6	\
284802	10	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	
284803	11	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	
284804	12	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	

284805	12	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708
284806	16	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617

	V7	V8	V9	...	V21	V22	V23	\
284802	-4.918215	7.305334	1.914428	...	0.213454	0.111864	1.014480	
284803	0.024330	0.294869	0.584800	...	0.214205	0.924384	0.012463	
284804	-0.296827	0.708417	0.432454	...	0.232045	0.578229	-0.037501	
284805	-0.686180	0.679145	0.392087	...	0.265245	0.800049	-0.163298	
284806	1.577006	-0.414650	0.486180	...	0.261057	0.643078	0.376777	

	V24	V25	V26	V27	V28	Amount	Class
284802	-0.509348	1.436807	0.250034	0.943651	0.823731	0.77	0
284803	-1.016226	-0.606624	-0.395255	0.068472	-0.053527	24.79	0
284804	0.640134	0.265745	-0.087371	0.004455	-0.026561	67.88	0
284805	0.123205	-0.569159	0.546668	0.108821	0.104533	10.00	0
284806	0.008797	-0.473649	-0.818267	-0.002415	0.013649	217.00	0

[5 rows x 31 columns]

```
[4]: Count_0 = df[df['Class']==0].groupby('Time').count()['Class']
      Count_1 = df[df['Class']==1].groupby('Time').count()['Class']
      df_counts = pd.concat([Count_0, Count_1], axis = 1, keys = ['Class 0', 'Class_
      ↪1'])
      df_counts = df_counts.reset_index()
      df_counts
```

```
[4]:
```

	Time	Class 0	Class 1
0	0	11655	9
1	1	11650	17
2	2	11799	17
3	3	11806	23
4	4	11960	24
5	5	11887	17
6	6	11870	23
7	7	11709	18
8	8	11806	19
9	9	11944	22
10	10	11865	16
11	11	11640	17
12	12	12015	19
13	13	11985	25
14	14	11893	21
15	15	11817	27
16	16	11779	24
17	17	11746	17
18	18	11723	23
19	19	11999	26

20	20	11908	22
21	21	12062	18
22	22	11911	28
23	23	11886	20


```
[5]: plt.plot(df_counts['Time'], df_counts['Class 0'])
plt.title('Non Fraud')
plt.xlabel('Hour')
plt.ylabel('Transactions')
plt.show()
```



```
[27]: plt.plot(df_counts['Time'],df_counts['Class 1'], color = 'red')
plt.title('Fraud')
plt.xlabel('Hour')
plt.ylabel('Transactions')
plt.show()
```



```
[17]: plt.figure(figsize = (8,6))  
plt.title('Original Dataset Correlation Credicard Transactions')  
sns.heatmap(df.corr(), linewidths = 1 ,cmap = 'Reds')  
plt.show()
```



2.4 Resampling Data

```
[10]: Class_0 = df[df['Class'] == 0]
sample_size = 492
sample_Class_0 = Class_0.sample(sample_size, random_state = 0)
print(sample_Class_0.value_counts(sample_Class_0['Class']))
print(sample_Class_0)
```

```
Class
0    492
dtype: int64
```

	Time	V1	V2	V3	V4	V5	V6 \
266085	15	2.049094	0.186189	-1.707198	0.530768	0.160589	-1.448570
172120	7	2.125540	-0.030714	-1.527653	0.121046	0.543172	-0.347988
15136	12	-4.155859	-5.705748	0.274699	-0.993262	-6.059393	5.210848
96393	15	-0.566420	-0.579576	0.823503	-1.451240	-0.583587	0.206381
208225	9	0.060858	-0.261762	-1.699493	-1.202327	3.699527	3.196249
...

170744	20	-0.995867	1.201192	0.395282	-0.886059	1.706389	-0.367065
222921	19	-1.008598	-0.075940	2.004425	-0.505601	0.509548	0.171588
188275	3	-1.773275	-1.718423	1.455868	-2.174910	1.570504	-0.734998
44186	2	-0.524514	1.353868	0.217128	1.241104	-0.045413	-0.880130
147585	8	-1.731798	1.418551	0.565101	1.196712	0.365576	-0.299215

	V7	V8	V9	...	V21	V22	V23	\
266085	0.239310	-0.353611	0.634425	...	0.197782	0.741141	-0.009744	
172120	0.157221	-0.229126	0.477999	...	-0.336497	-0.838932	0.275173	
15136	5.811316	0.367888	1.750710	...	1.371671	1.195815	4.188762	
96393	1.601392	-0.370446	-1.910354	...	-0.065082	-0.761357	0.641524	
208225	0.437208	0.421541	0.492435	...	0.008303	0.534602	0.089602	
...	
170744	1.089408	-0.210834	-0.608644	...	-0.379122	-1.111202	-0.469886	
222921	0.411154	0.066247	-0.024477	...	0.295670	0.738921	-0.229421	
188275	-0.974958	0.386686	-1.026933	...	0.360589	0.305160	-0.122747	
44186	0.197110	0.432705	-0.448191	...	0.133762	0.252190	0.073953	
147585	1.092355	-0.242769	-0.089824	...	-0.185723	-0.035647	-0.702781	

	V24	V25	V26	V27	V28	Amount	Class
266085	-0.085057	0.228384	-0.097292	-0.001028	-0.032390	2.99	0
172120	0.049145	-0.156765	0.205919	-0.072321	-0.059009	1.98	0
15136	-1.091077	1.033044	0.224493	-0.486741	0.194275	1937.66	0
96393	-0.568974	-0.053164	-0.690995	-0.228630	-0.157254	320.05	0
208225	0.667918	0.017798	0.611584	-0.469946	-0.514370	11.50	0
...
170744	-0.016865	0.621809	0.303837	-0.258417	0.056865	1.29	0
222921	0.770200	0.591492	-0.124258	-0.142162	-0.142230	82.86	0
188275	0.175912	0.810150	-0.155587	-0.073359	0.068604	34.70	0
44186	0.240549	-0.303189	-0.321691	-0.144621	-0.001791	1.00	0
147585	-0.237518	0.556394	-0.449458	-0.493690	-0.354016	108.00	0

[492 rows x 31 columns]

I created a new DataFrame for values with Class 0. Then I randomly choosed sample 492 elements, compared with Class 1 also 492 elements.Next, I will merge the new DataFrame with Class 0 and the DataFrame with Class 1 to create a balanced DataFrame. Let's check!

```
[11]: Class_1 = df[df['Class'] == 1]
df_balanced = pd.concat([Class_1, sample_Class_0], axis = 0, ignore_index =
↳ True)
df_balanced.head(5)
```

```
[11]:   Time    V1    V2    V3    V4    V5    V6    V7 \
0    22 -2.312227  1.951992 -1.609851  3.997906 -0.522188 -1.426545 -2.537387
1    16 -3.043541 -3.157307  1.088463  2.288644  1.359805 -1.064823  0.325574
2    22 -2.303350  1.759247 -0.359745  2.330243 -0.821628 -0.075788  0.562320
3     2 -4.397974  1.358367 -2.592844  2.679787 -1.128131 -1.706536 -3.496197
```

```
4      7  1.234235  3.019740 -4.304597  4.732795  3.624201 -1.357746  1.713445
```

```

      V8      V9  ...      V21      V22      V23      V24      V25  \
0  1.391657 -2.770089 ...  0.517232 -0.035049 -0.465211  0.320198  0.044519
1 -0.067794 -0.270953 ...  0.661696  0.435477  1.375966 -0.293803  0.279798
2 -0.399147 -0.238253 ... -0.294166 -0.932391  0.172726 -0.087330 -0.156114
3 -0.248778 -0.247768 ...  0.573574  0.176968 -0.436207 -0.053502  0.252405
4 -0.496358 -1.282858 ... -0.379068 -0.704181 -0.656805 -1.632653  1.488901

```

```

      V26      V27      V28  Amount  Class
0  0.177840  0.261145 -0.143276    0.00     1
1 -0.145362 -0.252773  0.035764   529.00     1
2 -0.542628  0.039566 -0.153029   239.93     1
3 -0.657488 -0.827136  0.849573    59.00     1
4  0.566797 -0.010016  0.146793    1.00     1

```


```
[5 rows x 31 columns]
```

```
[12]: new_class = df_balanced['Class'].value_counts()
      new_class
```


```
[12]: 1    492
      0    492
      Name: Class, dtype: int64
```

I successfully address the imbalanced data. Now, the result show that both values 0 and 1 have 492 instances. To make it easier to visualize, I will put them on a bar chart below. Let's check!

```
[13]: plt.bar(new_class.index, new_class)
      plt.title('Bar Chart for Class')
      plt.xlabel ('Class')
      plt.ylabel('Number of Transactions')
      plt.xticks([0, 1], ['Non Fraud (0)', 'Fraud (1)'])
      plt.show()
```



```
[14]: plt.figure(figsize = (8,6))  
plt.title('Subset Dataset Correlation of Credit Card Transactions')  
sns.heatmap(df_balanced.corr(),linewidths = 1, cmap = 'Blues')  
plt.show()
```

2.5 Data Preprocessing On Subset Dataset

In this part, I split the dataset into training and testing sets. X is features for classification variables to the target column y. It means, y contain 'Class' within there are 0 and 1. And the models will use the rest columns to classificate whether a variable is in 0 or is in 1.

```
[28]: X = df_balanced.drop(['Class'], axis = 1)
      y = df_balanced['Class']
      X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.
      ↪2,random_state = 10)
```

Next, I do standardizate data

```
[29]: sc = StandardScaler()
      X_train = sc.fit_transform(X_train)
      X_test = sc.transform(X_test)
```

2.6 Building Models

In this par, I will build 8 models machine learning to solve this case.

```
[3]: lr = LogisticRegression()
    dtc = DecisionTreeClassifier()
    rfc = RandomForestClassifier()
    gbc = GradientBoostingClassifier()
    knn = KNeighborsClassifier()
    svc = SVC()
    gnb = GaussianNB()
    xgb = XGBClassifier()
```

```
[37]: lr.fit(X_train, y_train)
    y_pred = lr.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    Classification_rp = classification_report(y_test, y_pred)
    print(Classification_rp)
    print(confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.84	0.96	0.90	95
1	0.96	0.83	0.89	102
accuracy			0.89	197
macro avg	0.90	0.90	0.89	197
weighted avg	0.90	0.89	0.89	197

```
[[91  4]
 [17 85]]
```

```
[39]: dtc.fit(X_train, y_train)
    y_pred = dtc.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    Classification_rp = classification_report(y_test, y_pred)
    print(Classification_rp)
    print(confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.87	0.93	0.90	95
1	0.93	0.87	0.90	102
accuracy			0.90	197
macro avg	0.90	0.90	0.90	197
weighted avg	0.90	0.90	0.90	197

```
[[88  7]
 [13 89]]
```

```
[40]: rfc.fit(X_train, y_train)
y_pred = rfc.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
Classification_rp = classification_report(y_test, y_pred)
print(Classification_rp)
print(confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.99	0.92	95
1	0.99	0.84	0.91	102
accuracy			0.91	197
macro avg	0.92	0.92	0.91	197
weighted avg	0.92	0.91	0.91	197

```
[[94  1]
 [16 86]]
```

```
[41]: gbc.fit(X_train, y_train)
y_pred = gbc.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
Classification_rp = classification_report(y_test, y_pred)
print(Classification_rp)
print(confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.98	0.91	95
1	0.98	0.84	0.91	102
accuracy			0.91	197
macro avg	0.92	0.91	0.91	197
weighted avg	0.92	0.91	0.91	197

```
[[93  2]
 [16 86]]
```

```
[42]: knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
Classification_rp = classification_report(y_test, y_pred)
print(Classification_rp)
print(confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.86	0.96	0.91	95
1	0.96	0.85	0.90	102
accuracy			0.90	197
macro avg	0.91	0.91	0.90	197
weighted avg	0.91	0.90	0.90	197

```
[[91  4]
 [15 87]]
```

```
[43]: svc.fit(X_train, y_train)
y_pred = svc.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
Classification_rp = classification_report(y_test, y_pred)
print(Classification_rp)
print(confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.96	0.90	95
1	0.96	0.84	0.90	102
accuracy			0.90	197
macro avg	0.90	0.90	0.90	197
weighted avg	0.90	0.90	0.90	197

```
[[91  4]
 [16 86]]
```

```
[44]: gnb.fit(X_train, y_train)
y_pred = gnb.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
Classification_rp = classification_report(y_test, y_pred)
print(Classification_rp)
print(confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.84	0.97	0.90	95
1	0.97	0.83	0.89	102
accuracy			0.90	197
macro avg	0.90	0.90	0.90	197
weighted avg	0.91	0.90	0.90	197

```
[[92  3]
```

[17 85]]

```
[45]: xgb.fit(X_train, y_train)
y_pred = xgb.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
Classification_rp = classification_report(y_test, y_pred)
print(Classification_rp)
print(confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.88	0.96	0.91	95
1	0.96	0.87	0.91	102
accuracy			0.91	197
macro avg	0.92	0.92	0.91	197
weighted avg	0.92	0.91	0.91	197

[[91 4]
[13 89]]

The results indicate that the models perform quite well on this data subset. To avoid bias opinions, I have selected more 3 models with lower accuracy (Logistic Regression, SVC, Decision Tree Classifier) to compare with the top 3 models (Random Forest Classifier, Gradient Boosting Classifier, XGB Classifier) with the highest accuracy score of 91% to re-run on the original dataset in order to assess and compare the effectiveness of these 6 models.

2.7 Re-Test On Original Dataset

```
[4]: X = df.drop(['Class'], axis = 1)
y = df['Class']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
↳ random_state = 11)
```

```
[5]: sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

These are 3 top result models on data subset:

```
[6]: rfc.fit(X_train, y_train)
y_pred = rfc.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
Classification_rp = classification_report(y_test, y_pred)
print(Classification_rp)
print(confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.00	1.00	1.00	56854
1	0.98	0.77	0.86	108
accuracy				1.00 56962
macro avg	0.99	0.88	0.93	56962
weighted avg	1.00	1.00	1.00	56962

```
[[56852    2]
 [   25   83]]
```

```
[7]: gbc.fit(X_train, y_train)
y_pred = gbc.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
Classification_rp = classification_report(y_test, y_pred)
print(Classification_rp)
print(confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56854
1	0.77	0.80	0.79	108
accuracy				1.00 56962
macro avg	0.89	0.90	0.89	56962
weighted avg	1.00	1.00	1.00	56962

```
[[56829    25]
 [   22   86]]
```

```
[8]: xgb.fit(X_train, y_train)
y_pred = xgb.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
Classification_rp = classification_report(y_test, y_pred)
print(Classification_rp)
print(confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56854
1	0.98	0.79	0.87	108
accuracy				1.00 56962
macro avg	0.99	0.89	0.94	56962
weighted avg	1.00	1.00	1.00	56962

```
[[56852    2]
 [   23   85]]
```

These are 3 lower result models on data subset:

```
[9]: svc.fit(X_train, y_train)
y_pred = svc.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
Classification_rp = classification_report(y_test, y_pred)
print(Classification_rp)
print(confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56854
1	0.97	0.67	0.79	108
accuracy			1.00	56962
macro avg	0.99	0.83	0.90	56962
weighted avg	1.00	1.00	1.00	56962

```
[[56852    2]
 [   36   72]]
```

```
[10]: dtc.fit(X_train, y_train)
y_pred = dtc.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
Classification_rp = classification_report(y_test, y_pred)
print(Classification_rp)
print(confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56854
1	0.87	0.80	0.83	108
accuracy			1.00	56962
macro avg	0.93	0.90	0.92	56962
weighted avg	1.00	1.00	1.00	56962

```
[[56841   13]
 [   22   86]]
```

```
[11]: lr.fit(X_train, y_train)
y_pred = lr.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
Classification_rp = classification_report(y_test, y_pred)
print(Classification_rp)
print(confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

	0	1.00	1.00	1.00	56854
	1	0.82	0.68	0.74	108
accuracy				1.00	56962
macro avg		0.91	0.84	0.87	56962
weighted avg		1.00	1.00	1.00	56962

```

[[56838    16]
 [   35    73]]

```

We can observe that the three models with lower accuracy rates (Logistic Regression, SVC, Decision Tree) on the subset of the data are likely to yield poor results when predicting the value 1 ('Fraud'), as the prediction rates are quite low. This suggests that these three models are not effective on the original full dataset. Conversely, the three models with the best results (Random Forest Classifier, Gradient Boosting, XGB) exhibit relatively high prediction rates for the value 1 ('Fraud'). This demonstrates that on this dataset, these three models have performed exceptionally well in identifying cases of 'Fraud,' even though there is a considerable imbalance in the dataset.

2.8 In Conclusion

From the conclusion above, we are only considering the top-performing three models on the original full dataset (Random Forest Classifier, Gradient Boosting, XGB). We can observe that the most effective performance comes from XGB, as this model has accurately identified cases of 1 ('Fraud') the best among the three models. Therefore, XGB stands out as the best-performing model for addressing this particular scenario.